

Dating Stats Atlas — data licensing and publication review

Prepared for counsel · 12 September 2026 · Nathan, sole founder

What we are asking for

A written memo covering (a) whether each data source listed in Part B may be used as described, and on what conditions, and (b) the two analysis questions in Part C that we cannot answer ourselves — our exposure on personalised rankings that include a race/ethnicity filter, and whether the distinction we draw between publishing a rendered page and distributing data holds up.

A per-source table of **yes / no / yes-with-conditions**, plus required attribution or disclaimer language, is the most useful form for us: our build system has a typed licence field on every data source and refuses to ship if a source marked non-shippable reaches the published artifact, so your answers get encoded and enforced automatically.

We are **not** asking about corporate structure, employment, tax, or fundraising.

Attached: dated copies of each source's licence or terms page, saved on the date shown in each filename.

Part A — The product, in one page

What it is. A pre-launch website that estimates, for one person's stated preferences, how many compatible potential partners live in each US metropolitan area, how many other people are competing for them, and what living there costs — then ranks the metro areas accordingly. It is a statistical reference tool, not a dating app: there are no profiles, no messaging, and no other users to meet.

Where the numbers come from. Almost entirely from US federal statistical agencies, chiefly the Census Bureau's American Community Survey public-use microdata — anonymised individual survey records, published by the Bureau for exactly this kind of analysis, with no names, addresses, or identifying detail. We aggregate those records to metropolitan-area population estimates. Nothing we publish describes an identifiable person.

What we collect from visitors: nothing. At launch there are no accounts, no login and no personal information. A visitor's preferences live in the page URL and are never stored.

Aggregate only, with suppression. Every figure displayed is a population estimate for a metro area of at least 250,000 people. Where the underlying sample is too small to support an estimate — fewer than 100 survey respondents — we suppress the result and say so rather than publishing a number.

Build time versus serve time. A build process reads the source data and produces a fixed, immutable data file. The website reads only that file. Some sources are read **only during**

the build, used as a quality check, and never contribute any value to the published file. This separation is enforced by an automated assertion that fails the build, not by staff discipline.

No public API. The site renders pages. There is no data feed, no bulk download and no documented public endpoint — the ranking endpoint is same-origin and undocumented. This is a deliberate decision (see question 3).

Identity filters. A visitor may optionally filter by race and ethnicity, using the eight standard categories derived from Census variables. Our own guardrails: the filter is off by default; we publish no page or headline of the form “best cities to meet [group] women”; metro pages describe population composition as demography; and whenever the filter is on we display alongside it the observed rate at which people of that group partner outside it, computed from the same public data, so the filter reads as the visitor’s own preference rather than a boundary. A later release will add a religion estimate, modelled from a public survey and labelled as modelled with a wider stated uncertainty.

Status and intent. Pre-launch, no public URL, no revenue. Intended to become a commercial product; business model not yet decided.

Part B — Data sources

Group 1 — federal sources that will ship at launch. Our understanding is that these are US government works, usable commercially with attribution. Question 1 asks you to confirm that and to specify required attribution or disclaimer language.

Source	Publisher	Provides	How used	Reaches the site?
ACS 5-year public-use microdata (PUMS), 2020–2024	Census Bureau	The core population estimates	Aggregated into our data file	Yes, as aggregates
ACS 5-year summary tables	Census Bureau	Median rent, group-quarters share, calibration totals	Context features; accuracy checks	Yes
2020 Census DHC (tables P1, P18)	Census Bureau	Tract-level population and group-quarters counts	Geographic allocation weights	Indirectly, via weights
2020 Census tract-to-PUMA relationship file	Census Bureau	Geographic crosswalk	Geographic allocation	Indirectly
OMB Bulletin 23-01 metro	OMB	Which counties form each	Defines the geography	Yes, as names

delineation		metro area		
Regional Price Parities	Bureau of Economic Analysis	Cost-of-living deflator	A scored feature	Yes
County Business Patterns	Census Bureau	Venue and establishment counts by industry	A scored feature	Yes
Quarterly Census of Employment and Wages	Bureau of Labor Statistics	Cross-check on the above	Accuracy check	No — check only
Smart Location Database v3.0	EPA	Street-network walkability measures	A scored feature	Yes
Climate Normals 1991–2020	NOAA	Pleasant-day counts	A scored feature	Yes
IPEDS 2024–25	NCES / Dept. of Education	Student population by metro	A scored feature	Yes
Crime Data Explorer	FBI	Offence rates	Displayed on metro pages with the FBI's own comparability caveat; never part of any ranking	Yes, as context

Note for question 1: we obtain some of this through the Census Bureau's API (which has its own terms of service, attached) and some by downloading bulk files from the Bureau's public file server. We would like to know whether those two routes carry different obligations.

Group 2 — planned for a later release.

Source	Publisher	Licence as we read it	Intended use
Overture Places	Overture Maps Foundation	Permissive, attribution required	Venue counts
Foursquare OS Places	Foursquare	Apache 2.0, NOTICE must be preserved	Venue counts, merged with the above

Cooperative
Election Study

Harvard Dataverse

CC0 1.0 (public
domain dedication)

The religion
estimate

Group 3 — read during the build, never published. Both forbid redistribution; one states it is not for marketing purposes. We are not seeking permission from either. Our plan is to read them during the build, compare them against our own modelled religion estimates as a quality check that can fail the build, and write no value derived from them into the published file. Question 4 asks whether that satisfies their terms.

Source

2020 US Religion Census

American Values Atlas county estimates

Group 4 — considered and excluded, on our own reading of their terms, listed so you can tell us if we have been more conservative than necessary: Yelp Fusion (forbids caching beyond 24 hours and building a competing listings database), Walk Score (forbids storing scores without written approval), Pew Religious Landscape restricted file (academic, non-commercial), General Social Survey sensitive files (faculty-gated, mandatory destruction, no AI input).

Group 5 — wanted, but not used pending your advice (question 5): Zillow Research and Redfin Data Center housing series, and the Opportunity Insights Social Capital Atlas. All three are widely used commercially; none publishes an explicit commercial grant we could find. We have built without them.

Part C — Questions

1. **Federal sources.** Can you confirm there is no licensing obstacle to commercial use of each Group 1 source, and specify the attribution or disclaimer language each requires? We know HUD requires a “not endorsed or certified by HUD User” disclaimer if we use HUD data; are there equivalents we have missed? Do the Census API terms of service impose obligations beyond those attaching to the same data downloaded as bulk files?
2. **Personalised rankings with a race or ethnicity filter.** This is our most important question. A visitor can filter by race and ethnicity, and the site then ranks metro areas for them personally, using public Census data. What is our exposure — fair-housing advertising rules, state consumer-protection law, or anything else we have not thought of? Are the guardrails described in Part A adequate, and what would you add or require? We would rather change the design now than add a disclaimer later.
3. **Rendered pages versus data distribution.** We publish no API deliberately. Our reasoning is that a rendered page is a work *produced from* data, whereas an API *distributes* the data itself, and that several ambiguous sources are comfortable with the former and not the latter. Is that distinction sound as a matter of licence interpretation? Does it materially reduce our exposure, and would opening a public API later require this review to be redone?

4. **Build-time-only use.** Does the Group 3 arrangement — read during the build, used as a comparison, never written into the published file, enforced by an automated check — satisfy terms that forbid redistribution and reposting and state that the data are not for marketing purposes?
5. **Group 5 sources.** What would we need in order to use Zillow, Redfin or Opportunity Insights data? Would publishing only a derived index, with attribution and without their underlying values, be sufficient?
6. **Attribution mechanics for Group 2.** What exactly must we display for Overture and for Foursquare OS Places, and where on the site must it appear?
7. **Name and domain.** Can we use the name “Dating Stats Atlas” and register the matching domain? Any obvious conflicts?
8. **Site policies.** Given no accounts and no personal information collected at launch, what do we need by way of a privacy policy and terms of use? We would like to state plainly that we do not store visitors’ search preferences, and want to be sure that statement is accurate and safe to make.

Part D — Attachments

Saved copies of each source’s licence or terms page, with the retrieval date in each filename, in attachments/.